

Measuring How Much IoT Devices Upload via Traffic Analysis

Enze Liu, TJ Smith, Zesen Zhang
University of California, San Diego
{e7liu,tjs003,zez003}@eng.ucsd.edu

ABSTRACT

We examine the problem of whether Amazon is being honest with users in presenting the audio data it collects. Amazon Echo is Amazon's smart speaker that has an always-on microphone. It starts recording after hearing a certain wake word (Alexa, Echo, etc.) and then sends the recorded audio to its server for further processing. Amazon allows users to browse and delete all audio data sent to them. We suspect that the audio data made available to users might only be a subset of the audio collected and transmitted to the server, and our initial experiments suggest that this might be the case.

1 INTRODUCTION

Physical devices connected to the internet, also known as Internet-of-things (IoT), have gained popularity in recent years. Among numerous types of IoT devices, smart speakers with voice assistants, such as Amazon Echo and Google Home, are widely seen in users' home. These devices detect and respond to voice commands. They normally have always-on microphones, which theoretically start recording after hearing a wake word and then send the voice data to a server via internet for further processing [1].

This behavior has sparked many privacy concerns, including but not limited to what is being recorded, how the collected data is used and stored, and whether it is being protected well [4–6, 8, 9]. Much of previous work has been centered around protecting sensitive information from being leaked to adversaries [4–6]. As for what data is being collected, Amazon has made it possible for users to view, play and delete their voice data stored on its server [7]. However, our intuition is that Amazon may reveal to users only part of the voice data being collected by its smart speakers. Up until now, little has been done to verify whether Amazon is being honest with us or not—is Amazon making public exactly the audio data it collects?

Our preliminary experiments with Amazon Echo have found that the recorded part of the same command is not consistent across different trials. This stimulates us to question and examine the behavior of Amazon Echo and dig into what is exactly being collected by it. We plan to use network traffic measurement for our work, through which previous researchers have been able to infer users' activities [6], and establish network signatures for Amazon Echo [7]. If this approach does work, we will be able to answer a whole set of research questions with it.

Permission to make digital or hard copies of all or part of this work for personal or classroom use is granted without fee provided that copies are not made or distributed for profit or commercial advantage and that copies bear this notice and the full citation on the first page. Copyrights for components of this work owned by others than ACM must be honored. Abstracting with credit is permitted. To copy otherwise, or republish, to post on servers or to redistribute to lists, requires prior specific permission and/or a fee. Request permissions from permissions@acm.org.
Conference'17, July 2017, Washington, DC, USA
© 2019 Association for Computing Machinery.
ACM ISBN 978-x-xxxx-xxxx-x/YY/MM...\$15.00
<https://doi.org/10.1145/nnnnnnnn.nnnnnnn>

2 BACKGROUND

Amazon Echo is a smart speaker developed by Amazon that is connected to the internet [3]. It has an always-on microphone that starts recording automatically after hearing certain wake word (Alexa, Echo, etc.). Voice commands following the wake word are streamed to Amazon's cloud-based intelligent command handling program, known as Alexa, for further processing. Alexa will then try to respond to users' commands. Amazon also keeps copies of the voice commands and responses, together known as response cards [7]. These response cards have been made available to users by Amazon so that they can be viewed, played, and deleted [2].

3 EXPERIMENT

We ran some experiments on Alexa and using Wireshark to see whether Alexa recorded consumer's voice more than we supposed they do. We then compared the size of package of each experiment to give out our results.

3.1 Environment Set Up

In order to catch the package that was sent out by Alexa to Amazon cloud, we use set up a hotspot from our server and let Alexa connected to it. In the meanwhile, we play a voice and record the time when Alexa replied our command. Further, we use Wireshark to record all the package that going through the hotspot which is connected by Alexa.

We did the experiment overnight in the office, to make sure there were no voice generated by person and all the thing Alexa might record was the voice played by our code.

3.2 Baseline Experiment

In the baseline experiment, we kept Alexa in the office overnight without doing anything and catch all the package that sent out by Alexa.

In this part, we assume that Alexa will not transfer any voice message to the cloud as there were no voice except some noise generated by nature. Therefore, we could use the package that were caught up in this experiment to be the baseline, who were the "necessary" package that Alexa communicate with cloud server. And we assume them to be unsuspicious package, which including some handshake package and broadcasting package to let the Alexa confirm which server to use.

3.3 Music Experiment

In the music experiment part, we played a music named Home overnight and catch up the data transmitted by Alexa.

"Home" is a music without any voice from human but it did contain some sound that could be recorded by Alexa. As Alexa has mentioned in its document[1]: "By default, Echo devices are designed to detect only your chosen wake word (Alexa, Amazon,

Computer or Echo)", it would not record and transmit anything to cloud once you did not call them. We would assume that Alexa would not transmit any voice package to cloud server as there is no human voice in the music, nothing would be recognized by the echo as "Alexa". As a result, if Alexa did the thing in the way as they announced, the package transmission result would be the same as it did in the Baseline Experiment.

3.4 "Prefix" Experiment

In this part, we did two different kinds of experiments. First, we played a voice with a 9 seconds prefix sentence and then followed by a 2.5 seconds command to wake Alexa up and let it transmit the package and answer. The prefix sentence is a random sentence that simulates a person talking in his daily life, and then he asked Alexa a question. The whole prefix sentence is: "I think I am having some trouble hearing what people trying to say to me. I am not feeling particularly well now. I am also thinking switching a new job." And the command is: "Alexa, where is New York City." It is really common in our daily life that people just having some conversation and suddenly find something to deal with, therefore they send a command to Alexa.

Here, we let Alexa to start transmitting the voice to the server. As Amazon has mentioned in his document [1]: "By default, Alexa-enabled devices only stream audio to the cloud if the wake word is detected (or Alexa is activated by pressing a button)." it will not transmit any data before the wake up word. Our assumption is that, if Alexa truly did the transmission as it mention, the size of each transmission package should be the same, as there is only a short command at the end of the sentence. However, if the optimization algorithm of Alexa is not efficient enough, it might transmit some parts of the prefix sentence to the cloud and the package size would wave whole night.

In order to make our results more significant, we did another comparison experiment. We played a voice with full prefix sentence mentioned above and a voice with no prefix sentence but just command for Alexa. And if Alexa did try some part of the prefix sentence, the package size of these two parts would show up a significant difference.

3.5 "Postfix" Experiment

Moreover, we want to measure whether Alexa will keep transmitting package even if Alexa has noticed the whole command and started replying. The assumption here is that, as Alexa has started replying, it means Alexa has fully recognized the whole command and it should not transmit the rest of the sentence onto the server. It fits our real life situation that someone sends a command to Alexa and starts another conversation with others immediately. If Alexa transmits the rest of the conversation, it would definitely rises the privacy concern.

Here, we use "Alexa, where is New York City" as a command sentence and adds a 1 second postfix sentence. We gave an approximately 0.5 seconds gap between the command and postfix sentence to make sure that Alexa has recognized the whole command and started to reply. Then we started playing three different kinds of sentences which are: a) command voice b) command voice with

0.5 second postfix sentence and c) command voice with 1 second postfix sentence. And analyze the package sent out by Alexa.

What's more, we want to see whether that 0.5 second pause is necessary for Alexa to realize the command is over and play the command with postfix sentence immediately. We try to find out whether Alexa would automatically cut off the voice once it recognized the following sentence is meaningless. We played command with 0,1,2,3,4,5,6 seconds postfix sentence, which contains nonsense sentences, to see whether Alexa would cut off the transmission itself.

Further, we try to figure out what is the long enough pause for Alexa to cut off the voice transmission and recognize the whole sentence. We played the command with different pause before playing the whole postfix sentence, to figure out when Alexa would reply a right answer to our command. Here we tried the voice gap between command and postfix sentence with 0.2,0.3,0.4,0.5 seconds.

REFERENCES

- [1] Amazon Echo (2nd generation) — Alexa Speaker. <https://www.amazon.com/Generation-improved-sound-powered-design/dp/B06XCM9LJ4?> (Accessed on 10/15/2019).
- [2] 2010. GP. (2010). <https://www.amazon.com/gp/help/customer/display.html?nodeId=201602040>
- [3] 2019. Amazon Echo. (Oct 2019). https://en.wikipedia.org/wiki/Amazon_Echo
- [4] Noah Aporthe, Danny Yuxing Huang, Dillon Reisman, Arvind Narayanan, and Nick Feamster. 2019. Keeping the smart home private with smart (er) iot traffic shaping. *Proceedings on Privacy Enhancing Technologies* 2019, 3 (2019), 128–148.
- [5] Noah Aporthe, Dillon Reisman, and Nick Feamster. 2017. A smart home is no castle: Privacy vulnerabilities of encrypted iot traffic. *arXiv preprint arXiv:1705.06805* (2017).
- [6] Noah Aporthe, Dillon Reisman, Srikanth Sundaresan, Arvind Narayanan, and Nick Feamster. 2017. Spying on the smart home: Privacy attacks and defenses on encrypted iot traffic. *arXiv preprint arXiv:1708.05044* (2017).
- [7] Marcia Ford and William Palmer. 2019. Alexa, are you listening to me? An analysis of Alexa voice service network traffic. *Personal and Ubiquitous Computing* 23, 1 (2019), 67–79.
- [8] Geoffrey Fowler. 2019. Perspective | Alexa has been eavesdropping on you this whole time. (May 2019). <https://www.washingtonpost.com/technology/2019/05/06/alexa-has-been-eavesdropping-you-this-whole-time/>
- [9] Josephine Lau, Benjamin Zimmerman, and Florian Schaub. 2018. Alexa, are you listening?: Privacy perceptions, concerns and privacy-seeking behaviors with smart speakers. *Proceedings of the ACM on Human-Computer Interaction* 2, CSCW (2018), 102.